

Ex3 - Handwritten character recognition using neural networks

Colab Link:

[Colab Link](#)

Aim:

To understand neural networks and perform the task of handwritten character recognition

Code & Output:

Import Dependencies

```
import tensorflow as tf
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import os
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D
from tensorflow.keras.layers import MaxPool2D
from tensorflow.keras.layers import Flatten
from tensorflow.keras.layers import Dense
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
import cv2
import time
import urllib.request
from PIL import Image
```

Reading And Sampling Data

```
df = pd.read_csv('english.csv')
df.head()
```

	image	label
0	Img/img001-001.png	0
1	Img/img001-002.png	0
2	Img/img001-003.png	0
3	Img/img001-004.png	0
4	Img/img001-005.png	0

df.columns

```
Index(['image', 'label'], dtype='object')
```

df.index

```
RangeIndex(start=0, stop=3410, step=1)
```

```
im = Image.open(df.iloc[np.random.randint(0,3410),0])  
im.show()
```

Pre Processing

Handling NULL Values

df['image'].isnull().sum()

```
0
```

df['label'].isnull().sum()

```
0
```

URL Image Read Function with OpenCV

```
def read_url_pic(x):  
    image_url = x  
  
    with urllib.request.urlopen(image_url) as url:  
        s = url.read()  
  
    arr = np.asarray(bytearray(s), dtype=np.uint8)  
    image = cv2.imdecode(arr, -1)  
    return image
```

Plot Image from URL with Matplotlib

```
def plot_url_pic(x):  
    image_url = x  
  
    with urllib.request.urlopen(image_url) as url:  
        s = url.read()  
  
    arr = np.asarray(bytearray(s), dtype=np.uint8)  
    image = cv2.imdecode(arr, -1)  
  
    plt.imshow(cv2.cvtColor(image, cv2.COLOR_BGR2RGB))  
    plt.show()
```

Convert Image to Array with OpenCV

```
def img_to_arr(x):  
    img = cv2.imread(x)  
    img = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)  
    img = cv2.resize(img, (64,64))  
    return img
```


Machine Learning Lab
UCS2612
Handwritten character recognition using neural networks

Ex. No: 3
12-2-24

Y.V.Ojus
3122 21 5001 125

Displaying Image Array Shape and Visualization

```
print(img_to_arr(df['image'])[0]).shape  
plt.imshow(img_to_arr(df['image'])[0]))
```



Parallel Image Processing with Multiprocessing

```
start_time = time.time()  
import multiprocessing  
  
with multiprocessing.Pool(4) as p:  
    images = p.map(img_to_arr, df['image'])  
end_time = time.time()  
print("with multiprocessing:", end_time-start_time)
```

```
with multiprocessing: 229.7383852005005
```


Reshape Image Array for Deep Learning Model Input

```
x = np.array(images).reshape(3410, 64, 64,1)  
x.shape
```

```
(3410, 64, 64, 1)
```

Visualizing Image with Corresponding Label

```
img = x[119].reshape((64,64))  
plt.imshow(img,cmap='gray')  
print(df['label'].iloc[119])
```



Label Encoding for Classification

```
y = df['label']  
le = LabelEncoder()  
y_label = le.fit_transform(y)
```

Splitting Data into Training and Testing Sets

```
train_images,test_images,train_labels,test_labels =  
train_test_split(x,y_label,test_size=0.2,random_state=42)  
train_images = train_images/255.0  
test_images = test_images/255.0
```

Convolutional Neural Network Model Training

```
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout,  
BatchNormalization
```

```
model = Sequential()
```

```
model.add(Conv2D(512, (5, 5), activation='relu', input_shape=(64, 64, 1)))
```

```
model.add(MaxPooling2D(2, 2))
```

```
model.add(Conv2D(256, (3, 3), activation='relu'))
```

```
model.add(MaxPooling2D(2, 2))
```

```
model.add(Conv2D(256, (3, 3), activation='relu'))
```

```
model.add(MaxPooling2D(2, 2))
```

```
model.add(Flatten())
```

```
model.add(Dense(256, activation='relu'))
```

```
model.add(Dropout(0.5))
```

```
model.add(Dense(128, activation='relu'))
```

```
model.add(Dropout(0.5))
```

```
model.add(Dense(62, activation='softmax'))
```

```
Epoch 1/20  
86/86 [=====] - 6s 41ms/step - loss: 4.1317 - accuracy: 0.0172 - val_loss: 4.1273 - val_accuracy: 0.0191  
Epoch 2/20  
86/86 [=====] - 3s 37ms/step - loss: 4.0511 - accuracy: 0.0330 - val_loss: 3.8516 - val_accuracy: 0.0777  
Epoch 3/20  
86/86 [=====] - 3s 40ms/step - loss: 3.7260 - accuracy: 0.0806 - val_loss: 3.2224 - val_accuracy: 0.1921  
Epoch 4/20  
86/86 [=====] - 4s 41ms/step - loss: 3.2124 - accuracy: 0.1675 - val_loss: 2.5603 - val_accuracy: 0.3666  
Epoch 5/20  
86/86 [=====] - 3s 39ms/step - loss: 2.5948 - accuracy: 0.3050 - val_loss: 1.8749 - val_accuracy: 0.4897  
Epoch 6/20  
86/86 [=====] - 3s 38ms/step - loss: 2.1333 - accuracy: 0.4054 - val_loss: 1.4942 - val_accuracy: 0.5909  
Epoch 7/20  
86/86 [=====] - 3s 40ms/step - loss: 1.7806 - accuracy: 0.4703 - val_loss: 1.1092 - val_accuracy: 0.6525  
Epoch 8/20  
86/86 [=====] - 3s 37ms/step - loss: 1.5071 - accuracy: 0.5403 - val_loss: 0.9838 - val_accuracy: 0.7199  
Epoch 9/20  
86/86 [=====] - 3s 37ms/step - loss: 1.2936 - accuracy: 0.6008 - val_loss: 0.9504 - val_accuracy: 0.7111  
Epoch 10/20  
86/86 [=====] - 3s 39ms/step - loss: 1.1632 - accuracy: 0.6499 - val_loss: 0.8682 - val_accuracy: 0.7419  
Epoch 11/20  
86/86 [=====] - 3s 39ms/step - loss: 1.0106 - accuracy: 0.6859 - val_loss: 0.7958 - val_accuracy: 0.7639  
Epoch 12/20  
86/86 [=====] - 3s 38ms/step - loss: 0.8861 - accuracy: 0.7137 - val_loss: 0.7585 - val_accuracy: 0.7713  
Epoch 13/20  
86/86 [=====] - 3s 38ms/step - loss: 0.8085 - accuracy: 0.7302 - val_loss: 0.7352 - val_accuracy: 0.7727  
Epoch 14/20  
86/86 [=====] - 3s 38ms/step - loss: 0.7586 - accuracy: 0.7588 - val_loss: 0.7822 - val_accuracy: 0.7625  
Epoch 15/20  
86/86 [=====] - 3s 39ms/step - loss: 0.7127 - accuracy: 0.7625 - val_loss: 0.7015 - val_accuracy: 0.7845  
Epoch 16/20  
86/86 [=====] - 3s 38ms/step - loss: 0.5929 - accuracy: 0.8002 - val_loss: 0.6656 - val_accuracy: 0.7977  
Epoch 17/20  
86/86 [=====] - 3s 37ms/step - loss: 0.5764 - accuracy: 0.8109 - val_loss: 0.6602 - val_accuracy: 0.7962  
Epoch 18/20  
86/86 [=====] - 3s 37ms/step - loss: 0.5377 - accuracy: 0.8196 - val_loss: 0.6942 - val_accuracy: 0.7918  
Epoch 19/20  
86/86 [=====] - 3s 39ms/step - loss: 0.5230 - accuracy: 0.8284 - val_loss: 0.6490 - val_accuracy: 0.8006  
Epoch 20/20  
86/86 [=====] - 3s 38ms/step - loss: 0.4320 - accuracy: 0.8614 - val_loss: 0.6829 - val_accuracy: 0.7991  
<keras.src.callbacks.History at 0x7bc9a04ad030>
```

Display Model Summary

model.summary()

```
Model: "sequential_3"
```

Layer (type)	Output Shape	Param #
conv2d_7 (Conv2D)	(None, 60, 60, 512)	13312
max_pooling2d_6 (MaxPooling2D)	(None, 30, 30, 512)	0
conv2d_8 (Conv2D)	(None, 28, 28, 256)	1179904
max_pooling2d_7 (MaxPooling2D)	(None, 14, 14, 256)	0
conv2d_9 (Conv2D)	(None, 12, 12, 256)	590080
max_pooling2d_8 (MaxPooling2D)	(None, 6, 6, 256)	0
flatten_2 (Flatten)	(None, 9216)	0
dense_6 (Dense)	(None, 256)	2359552
dropout (Dropout)	(None, 256)	0
dense_7 (Dense)	(None, 128)	32896
dropout_1 (Dropout)	(None, 128)	0
dense_8 (Dense)	(None, 62)	7998

```
=====  
Total params: 4183742 (15.96 MB)  
Trainable params: 4183742 (15.96 MB)  
Non-trainable params: 0 (0.00 Byte)
```

Inverse Transform Predicted and Actual Labels

```
predicted_labels = le.inverse_transform(model.predict(test_images).argmax(axis=1))  
actual_labels = le.inverse_transform(test_labels)
```

```
22/22 [=====] - 1s 13ms/step
```

Display Predicted and Actual Labels for a Specific Index


```
print('predicted label is ',predicted_labels[67])  
print('Actual Label is:',actual_labels[67])
```

```
predicted label is ; j  
Actual Label is : j
```

Displaying Predicted and Actual Labels for a Specific Image

```
# number of image to be predicted  
# change the value of i  
i = 450
```

```
plt.imshow(test_images[i],cmap='gray')  
print('printed label:',predicted_labels[i])  
print('actual label:', actual_labels[i])
```



Evaluate Model Performance with Various Metrics

```
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score,
confusion_matrix
```

```
y_pred = model.predict(test_images)
y_pred_labels = np.argmax(y_pred, axis=1)
```

```
# Calculate evaluation metrics
accuracy = accuracy_score(test_labels, y_pred_labels)
precision = precision_score(test_labels, y_pred_labels, average='weighted')
recall = recall_score(test_labels, y_pred_labels, average='weighted')
f1 = f1_score(test_labels, y_pred_labels, average='weighted')
confusion_mat = confusion_matrix(test_labels, y_pred_labels)
```

```
# Print or use the evaluation metrics as needed
print("Accuracy:", accuracy)
print("Precision:", precision)
print("Recall:", recall)
print("F1 Score:", f1)
print("Confusion Matrix:")
print(confusion_mat)
```

```
22/22 [=====] - 0s 12ms/step
Accuracy: 0.7991202346041055
Precision: 0.8160896386154118
Recall: 0.7991202346041055
F1 Score: 0.7966589813430557
Confusion Matrix:
[[ 8  0  0 ...  0  0  0]
 [ 0  5  0 ...  0  0  0]
 [ 0  0  8 ...  0  0  0]
 ...
 [ 0  0  0 ... 12  0  2]
 [ 0  0  0 ...  0  8  0]
 [ 0  0  0 ...  0  0  9]]
```